

Exercise 1

Deadline 6.11.2025

Anna Bütfering

Lea Neumann

Juan Provencio

1. Two dimensional data

```
In [109... from sklearn.datasets import make_moons
import numpy as np
import matplotlib.pyplot as plt

from numpy import exp, sqrt, pi
```

```
In [110... # Plot
X, y = make_moons(n_samples = 500, noise = 0.1, random_state = 42)
plt.scatter(X[:, 0], X[:, 1], s=10)

# Implement a class for each of the following models, including
# methods for fit(), sample(), and score_samples()
```

Next, we will implement classes, à la sklearn, providing methods for `fit()`, `sample()` and `score_samples()` without using the pre-defined classes, but formulating them independently.

We will create a class for each of the following models:

1. two-dimensional histogram
2. single Gaussian
3. Gaussian mixture model
4. kernel density estimator with a Gaussian kernel

For evaluation, we will implement the maximum mean discrepancy metric with squared exponential and inverse multi-quadratic kernels

$$\text{se}(x, x') = \exp\left(-\frac{\|x - x'\|^2}{h^2}\right)$$

$$\text{imq}(x, x') = \frac{1}{\sqrt{1 + \frac{\|x - x'\|^2}{h^2}}}$$

1.1 Two dimensional histogram

```
In [111]... class two_dim_histogram:
    def __init__(self, random_state=None, bin_size=None):
        self.densities = None
        self.bin_edges_1 = None
        self.bin_edges_2 = None
        self.samples = None
        self.score = None
        self.log_density = None
        self.bin_indices = None
        self.bin_size = bin_size
        self.MMD_sq = None
        self.random_state = np.random.RandomState(random_state) if random_st
# Method to fit 2d histogram
'''
As in 1d, need to determine p_k and bin sizes.

In 1d, determined the bin sizes with the interquantile range IQR

    h_1D = 2 IQR / (N)^(1/3)

In 2d, do the same in both dimensions
'''
    def fit(self, X): # training_data := X
        # Get n_samples
        n_samples = X.shape[0]

        # Determine Interquantile range
```

```

# IQR = x_[3N/4] - x_[N/4]
"""
DONT WANT TO GIVE PEOPLE THE OPTION TO PICK THEIR OWN BINS BECAUSE
IM TOO TIRED
UGHH OK I NEED TO
"""
if self.bin_size == None:
    IQR_x1 = np.percentile(X[:, 0], 75) - np.percentile(X[:, 0], 25)
    IQR_x2 = np.percentile(X[:, 1], 75) - np.percentile(X[:, 1], 25)

    # Bin widths
    h_1 = 2 * IQR_x1 / ((n_samples) ** (1/3))
    h_2 = 2 * IQR_x2 / ((n_samples) ** (1/3))
    self.bin_size = (np.round(h_1, 2), np.round(h_2, 2))

    # Create bins k
    x1_min, x1_max = np.min(X[:, 0]), np.max(X[:, 0])
    x2_min, x2_max = np.min(X[:, 1]), np.max(X[:, 1])
    self.bin_edges_1 = np.arange(x1_min, x1_max + self.bin_size[0],
    self.bin_edges_2 = np.arange(x2_min, x2_max + self.bin_size[1],

else:
    x1_min, x1_max = np.min(X[:, 0]), np.max(X[:, 0])
    x2_min, x2_max = np.min(X[:, 1]), np.max(X[:, 1])
    self.bin_edges_1 = np.arange(x1_min, x1_max + self.bin_size[0],
    self.bin_edges_2 = np.arange(x2_min, x2_max + self.bin_size[1],

# Determine counts in bins k
counts, _, _ = np.histogram2d(X[:, 0], X[:, 1], bins = [self.bin_edges_1,
self.densities = counts / counts.sum()

return self.densities, self.bin_edges_1, self.bin_edges_2

# Sample
def sample(self, n_samples=1):
    # now we can draw samples from the learned histogram (generated new
    probs_flat = self.densities.ravel() # flatten the 2D probs array
    self.bin_indices = self.random_state.choice(probs_flat.size, size=n_
    ix, iy = np.unravel_index(self.bin_indices, self.densities.shape) #
    x0 = self.bin_edges_1[ix]
    x1 = self.bin_edges_1[ix + 1]
    y0 = self.bin_edges_2[iy]
    y1 = self.bin_edges_2[iy + 1]

    x_samples = x0 + self.random_state.rand(n_samples) * (x1 - x0)
    y_samples = y0 + self.random_state.rand(n_samples) * (y1 - y0)

    self.samples = np.column_stack([x_samples, y_samples]) # shape (n_s

    return self.samples

# Method to score samples
def score_samples(self, X):
    """Compute the log density for each point in X."""
    if self.densities is None or self.bin_edges_1 is None or self.bin_ec
        raise ValueError("Model is not fitted. Call the `fit` method fir

```

```

# Find the bin indices for each point in X
bin_indices_x = np.digitize(X[:, 0], self.bin_edges_1) - 1
bin_indices_y = np.digitize(X[:, 1], self.bin_edges_2) - 1

# Ensure indices are within range
valid_mask = (
    (bin_indices_x >= 0) & (bin_indices_x < self.densities.shape[0])
    & (bin_indices_y >= 0) & (bin_indices_y < self.densities.shape[1])
)

# Initialize log densities with -inf for points outside the histogram
self.log_densities = np.full(X.shape[0], -np.inf)

# Assign log densities for valid points
self.log_densities[valid_mask] = np.log(
    self.densities[bin_indices_x[valid_mask], bin_indices_y[valid_mask]]
)

return self.log_densities

# Compute mmd^2
def compute_mmd(self, X, h = 1, kernel = "squared_exponential"):
    if self.samples is None:
        raise ValueError("Samples are not generated. Call the `sample` method")
    # Implement MMD^2 metric with se, imq kernels
    # Need to vectorize MMD calculation because for loops (sums) are very slow
    # MMD^2 = Term1 + Term2 - Term3, can neglect Term1 because it is much smaller
    MMD^2 = 0
    for i in range(1, N):
        for i_prime in range(1, N_prime):
            MMD^2 += k(x_i, x_i_prime)

    for j in range(1, N_prime):
        for j_prime in range(1, N_prime):
            MMD^2 += k(x_prime_j, x_prime_j_prime)

    for i in range(1, N):
        for j in range(1, N_prime):
            MMD^2 += -2 k(x_i, x_prime_j)
    """
    # Instead, we want to calculate it using something like this:
    """
    MMD^2 = k(distance_matrix_x).sum()
        + k(distance_matrix_x_prime).sum()
        - 2 k(distance_matrix_xx_prime).sum()
    """
    # Calculate Distance matrices D as:
    # D[i, j] = || x_i - x_i_prime ||_2^2 = ||x||_2^2 + ||x_prime||_2^2 - 2 x · x_prime
    # D_XX = self.calculate_distance_matrix(X, X) Can be neglected
    D_XprimeXprime = self.calculate_distance_matrix(self.samples, self.samples)
    D_XXprime = self.calculate_distance_matrix(X, self.samples)

    # Depending on selected kernel, calculate MMD^2

```

```

N = X.shape[0]
Nprime = self.samples.shape[0]
if kernel == "squared_exponential":
    self.MMD_sq = (self.squared_exponential(D_XprimeXprime, h).sum()
                  - 2 * self.squared_exponential(D_XXprime, h).sum()) / (
elif kernel == "inverse_multi_quadratic":
    self.MMD_sq = (self.inverse_multi_quadratic(D_XprimeXprime, h).sum()
                  - 2 * self.inverse_multi_quadratic(D_XXprime, h).sum())
else:
    raise ValueError("Please select either 'squared_exponential' or
    """

    In the lecture we derived the formula where we sum over i, then sum
    k(i, i) are not 0, but 1. We still sum over them in the vectorization
    we adjust the denominator to N^2, N'^2 instead of N(N-1), N'(N'-1)
    """

    self.MMD_sq

    return self.MMD_sq

def calculate_distance_matrix(self, A, B):
    # D_ij = ||A_i - B_j||^2 = ||A_i||^2 + ||B_j||^2 - 2 A_i @ B_j.T
    # A_sq in (N_A, 1)
    A_sq = np.sum(A ** 2, axis = 1, keepdims = True)
    # B_sq in (1, N_B)
    B_sq = np.sum(B ** 2, axis = 1, keepdims = True).T # .T otherwise we
    AB_T = A @ B.T
    D_AB = A_sq + B_sq - 2 * AB_T
    return D_AB

def squared_exponential(self, D_AB, h):
    return exp(- D_AB / (h ** 2))

def inverse_multi_quadratic(self, D_AB, h):
    return 1 / sqrt(1 + D_AB / (h ** 2))

# Fit 2d histogram model
model = two_dim_histogram()
model.fit(X)
model.sample(2000)
model.score_samples(X)


make_moons_histogram2d, edges_1, edges_2 = model.fit(X)

# Sample from fitted model
X_samples = model.sample(n_samples = 2000)

# Plot
plt.figure()
plt.title("2D Histogram Model for 2 Moons Dataset")
plt.scatter(X[:, 0], X[:, 1], s=10, label = "Original Data")
plt.scatter(X_samples[:, 0], X_samples[:, 1], s=10, alpha = 0.4, label = "Sampled Data")
plt.legend()
plt.show()

# Score

```



1.2 Single Gaussian

```
In [112... from scipy.stats import multivariate_normal

class single_gaussian:
    def __init__(self):
        self.mu = None
        self.Sigma = None
        self.samples = None
        self.score = None
        self.MMD_sq = None
        self.log_density = None
    # Method to fit a single Gaussian
    """
    Need to determine appropriate mu, sigma for a single Gaussian on top of
    """
    def fit(self, X): # training_data := X
        # Gaussian with simply mu = np.mean, Sigma = np.cov
        self.mu = np.mean(X, axis = 0)
        self.Sigma = np.cov(X, rowvar = False)

        return self.mu, self.Sigma

    def sample(self, n_samples):
        self.samples = np.random.multivariate_normal(self.mu, self.Sigma, n_samples)
        return self.samples

    def score_samples(self, X):
```

```

"""Compute the log density for each point in X."""
if self.mu is None or self.Sigma is None:
    raise ValueError("Model parameters are not initialized. Call `fit` before using.")

# Compute the log density using the multivariate normal distribution
self.log_density = multivariate_normal.logpdf(X, mean=self.mu, cov=self.Sigma)
return self.log_density

def compute_mmd(self, X, h = 1, kernel = "squared_exponential"):
    # Implement MMD^2 metric with se, imq kernels
    # Need to vectorize MMD calculation because for loops (sums) are very slow
    # MMD^2 = Term1 + Term2 - Term3, can neglect Term1 because it is mod invariant
    """
    MMD^2 = 0
    for i in range(1,N):
        for i_prime != i:
            MMD^2 += k(x_i, x_i_prime)

    for j in range(1, N_prime):
        for j_prime != j:
            MMD^2 += k(x_prime_j, x_prime_j_prime)

    for i in range(1, N):
        for j in range(1, N_prime):
            MMD^2 += -2 k(x_i, x_prime_j)
    """
    # Instead, we want to calculate it using something like this:
    """
    MMD^2 = k(distance_matrix_x).sum()
            + k(distance_matrix_x_prime).sum()
            - 2 k(distance_matrix_xx_prime).sum()
    """
    # Calculate Distance matrices D as:
    # D[i, j] = || x_i - x_i_prime ||_2^2 = ||x||_2^2 + ||x_prime||_2^2 - 2 x^T x_prime
    # D_XX = self.calculate_distance_matrix(X, X) Can be neglected
    D_XprimeXprime = self.calculate_distance_matrix(self.samples, self.samples)
    D_XXprime = self.calculate_distance_matrix(X, self.samples)

    # Depending on selected kernel, calculate MMD^2
    N = X.shape[0]
    Nprime = self.samples.shape[0]
    if kernel == "squared_exponential":
        self.MMD_sq = (self.squared_exponential(D_XprimeXprime, h).sum()
                       - 2 * self.squared_exponential(D_XXprime, h).sum()) / (N * Nprime)
    elif kernel == "inverse_multi_quadratic":
        self.MMD_sq = (self.inverse_multi_quadratic(D_XprimeXprime, h).sum()
                       - 2 * self.inverse_multi_quadratic(D_XXprime, h).sum()) / (N * Nprime)
    else:
        raise ValueError("Please select either 'squared_exponential' or 'inverse_multi_quadratic'")
    """
    In the lecture we derived the formula where we sum over i, then sum over j, but k(i, i) are not 0, but 1. We still sum over them in the vectorization. To avoid this, we adjust the denominator to N^2, N'^2 instead of N(N-1), N'(N'-1)
    """

    self.MMD_sq = (self.MMD_sq * N * Nprime) / (N**2 + Nprime**2)

```

```

        return self.MMD_sq

    def calculate_distance_matrix(self, A, B):
        #  $D_{ij} = ||A_i - B_j||^2 = ||A_i||^2 + ||B_j||^2 - 2 A_i @ B_j.T$ 
        # A_sq in (N_A, 1)
        A_sq = np.sum(A ** 2, axis = 1, keepdims = True)
        # B_sq in (1, N_B)
        B_sq = np.sum(B ** 2, axis = 1, keepdims = True).T # .T otherwise wc
        AB_T = A @ B.T
        D_AB = A_sq + B_sq - 2 * AB_T
        return D_AB

    def squared_exponential(self, D_AB, h):
        return exp(- D_AB / (h ** 2))


    def inverse_multi_quadratic(self, D_AB, h):
        return 1 / sqrt(1 + D_AB / (h ** 2))

# Fit 2d single Gaussian model
model = single_gaussian()
make_moons_single_gaussian = model.fit(X)

# Sample from Gaussian
samples = model.sample(n_samples = 1000)
plt.figure()
plt.title("Single Gaussian model for 2 Moons dataset")
plt.scatter(X[:, 0], X[:, 1], s=10,
            label = "Dataset")
plt.scatter(samples[:, 0], samples[:, 1],
            alpha = 0.5,
            s = 10,
            label = "Samples")
plt.legend()
plt.show()

# Score

```

1.3 Gaussian Mixture Model

```
In [113]: from scipy.stats import multivariate_normal

class gaussian_mixture_model:
    def __init__(self, n_terms = 1, steps = 100):
        self.mus = None
        self.Sigmas = None
        self.Pis = None
        self.samples = None
        self.score = None
        self.n_terms = n_terms
        self.steps = steps

    ...

    p_hat = sum_{m = 1}^M Pi_m N(mu_m, Sigma_m), where mu_m, Sigma_m
    can be found through expectation maximization algorithm from 1D.

    Basics:
    (0) Initialize Sigma_k = Id, mu_k random datapoint, Pi_k^(0) = 1 / M
    (1) for step t or until convergence:
        a) calc. performance
            gamma_ik^(t) = Pi_k^(t-1) N(x_i | mu_k^(t-1), Sigma_k^(t-1)) / p
        b) update mu_k^(t), Sigma_k^(t)
            Pi_k^(t) = 1/N sum_{i=1}^N gamma_ik
            mu_k^(t) = sum_{i=1}^N gamma_ik x_i / sum_{i=1}^N gamma_ik
            Sigma_k^(t) ...

    ...
```

```

def initialize_parameters(self, X):
    N, D = X.shape
    M = self.n_terms

    # Initialize for t = 0
    random_indices = np.random.choice(N, M, replace = False)
    self.mus = X[random_indices]
    self.Pis = np.ones(M) / M
    self.Sigmas = np.array([np.eye(D) for _ in range(M)])

# Method to fit GMM
def fit(self, X): # training_data := X

    # Initialize parameters
    self.initialize_parameters(X)
    self.N, self.D = X.shape
    self.M = self.n_terms

    # Expectation Maximization algorithm
    for t in range(self.steps):
        self.gamma_ik = np.zeros((self.N, self.M))
        # a) Calculate gamma_ik
        for k in range(self.M):
            self.gamma_ik[:, k] = self.Pis[k] * multivariate_normal.pdf(

        # Normalize gamma_ik
        self.gamma_ik /= np.sum(self.gamma_ik, axis = 1, keepdims = True)

        # b) Update parameters
        # Calculate mus
        self.Pis = np.mean(self.gamma_ik, axis = 0)
        self.mus = (self.gamma_ik.T @ X) / np.sum(self.gamma_ik, axis =
        # Calculate Sigmas: sum_i gamma_ik (x - mu)^T (x - mu) / sum_i g
        # Define a small regularization constant to keep cov positive de
        REG_COV = 1e-6
        for k in range(self.M):
            X_centered = X - self.mus[k]
            weighted_X_centered = X_centered * self.gamma_ik[:, k][:, np
            self.Sigmas[k] = weighted_X_centered.T @ X_centered / np.sum

    return self.mus, self.Sigmas

def sample(self, n_samples):
    # Sample from GMM
    # 1. Pick from which Gaussian to sample based on weighs Pi_k
    component_indices = np.random.choice(self.M, size = n_samples, p = s
    # Initialize samples array
    self.samples = np.zeros((n_samples, self.D))
    # 2. Sample from selected Gaussian
    for k in range(self.M):
        # Count instances where Gaussian k was selected
        n_k = np.sum(component_indices == k)
        if n_k > 0:
            self.samples[component_indices == k] = np.random.multivariat
            self.mus[k], self.Sigmas[k], n_k)

```

```

        return self.samples

# Method to score samples
def score_samples(self, X):
    """Compute the log density for each point in X."""
    if self.Pis is None or self.mus is None or self.Sigmas is None:
        raise ValueError("Model parameters are not initialized. Call `fit`")

    # Compute the weighted density for each component
    densities = np.zeros((X.shape[0], self.M))
    for k in range(self.M):
        densities[:, k] = self.Pis[k] * multivariate_normal.pdf(X, mean=self.mus[k], cov=self.Sigmas[k])

    # Sum the densities across components
    total_density = np.sum(densities, axis=1)

    # Return the log density
    log_density = np.log(total_density)
    return log_density

# Method to evaluate
def compute_mmd(self, X, h = 1, kernel = "squared_exponential"):
    # Implement MMD^2 metric with se, img kernels
    # Need to vectorize MMD calculation because for loops (sums) are very slow
    # MMD^2 = Term1 + Term2 - Term3, can neglect Term1 because it is much smaller than Term2
    """
    MMD^2 = 0
    for i in range(1,N):
        for i_prime != i:
            MMD^2 += k(x_i, x_i_prime)

    for j in range(1, N_prime):
        for j_prime != j:
            MMD^2 += k(x_prime_j, x_prime_j_prime)

    for i in range(1, N):
        for j in range(1, N_prime):
            MMD^2 += -2 k(x_i, x_prime_j)
    """
    # Instead, we want to calculate it using something like this:
    """
    MMD^2 = k(distance_matrix_x).sum()
            + k(distance_matrix_x_prime).sum()
            - 2 k(distance_matrix_xx_prime).sum()
    """
    # Calculate Distance matrices D as:
    # D[i, j] = || x_i - x_i_prime ||_2^2 = ||x||_2^2 + ||x_prime||_2^2 - 2 x · x_prime
    # D_XX = self.calculate_distance_matrix(X, X) Can be neglected
    D_XprimeXprime = self.calculate_distance_matrix(self.samples, self.samples)
    D_XXprime = self.calculate_distance_matrix(X, self.samples)

    # Depending on selected kernel, calculate MMD^2
    N = X.shape[0]
    Nprime = self.samples.shape[0]
    if kernel == "squared_exponential":

```

```

        MMD_sq = (self.squared_exponential(D_XprimeXprime, h).sum() / (N
            - 2 * self.squared_exponential(D_XXprime, h).sum() / (N
elif kernel == "inverse_multi_quadratic":
        MMD_sq = (self.inverse_multi_quadratic(D_XprimeXprime, h).sum()
            - 2 * self.inverse_multi_quadratic(D_XXprime, h).sum()
else:
    raise ValueError("Please select either 'squared_exponential' or
    """

    In the lecture we derived the formula where we sum over i, then sum
    k(i, i) are not 0, but 1. We still sum over them in the vectorization
    we adjust the denominator to N^2, N'^2 instead of N(N-1), N'(N'-1)
    """

    self.MMD_sq = MMD_sq

    return self.MMD_sq

def calculate_distance_matrix(self, A, B):
    # D_ij = ||A_i - B_j||^2 = ||A_i||^2 + ||B_j||^2 - 2 A_i @ B_j.T
    # A_sq in (N_A, 1)
    A_sq = np.sum(A ** 2, axis = 1, keepdims = True)
    # B_sq in (1, N_B)
    B_sq = np.sum(B ** 2, axis = 1, keepdims = True).T # .T otherwise wrong
    AB_T = A @ B.T
    D_AB = A_sq + B_sq - 2 * AB_T
    return D_AB

def squared_exponential(self, D_AB, h):
    return exp(- D_AB / (h ** 2))


def inverse_multi_quadratic(self, D_AB, h):
    return 1 / sqrt(1 + D_AB / (h ** 2))

# Test GMM
model = gaussian_mixture_model(n_terms = 10, steps = 100)
make_moons_gmm = model.fit(X)

# Plot calculated GMM over data
samples = model.sample(1000)
plt.figure()
plt.title("GMM model for 2 Moons dataset")
plt.scatter(X[:, 0], X[:, 1], s=10,
            label = "Dataset")
plt.scatter(samples[:, 0], samples[:, 1],
            alpha = 0.5,
            s = 10,
            label = "GMM Samples")
plt.legend()
plt.show()

# Score

```



Evaluation of Models

Steps:

1. Vary hyperparameters for each model (e.g., number of components for GMM).
2. Compute MMD for different bandwidths (h).
3. Visualize the results:
 - MMD values as a function of hyperparameters.
 - Density estimates and generated samples.

```
In [114... # Define data to train and test
X_train, _ = make_moons(n_samples = 1000, noise = 0.1, random_state = 42)
X_test, _ = make_moons(n_samples = 1000, random_state = 24)
```

2D Histogram

```
In [115... # Bin sizes
bin_sizes = [(0.01, 0.01), (0.02, 0.02), (0.03, 0.03), (0.05, 0.05), (0.1, 0.1)]
results = {}
avg_bin_sizes = []
# Make a grid of plots
fig, axs = plt.subplots(2, 4, figsize=(15, 10))
fig.suptitle("2D Histogram Models with Different Bin Sizes", fontsize=16)
for bin_size in bin_sizes:
    model = two_dim_histogram(bin_size=bin_size)
    model.fit(X_train)
```

```

model.sample(2000)
model.score_samples(X_test)
mmd_sq_se = model.compute_mmd(X_test, h = 0.5, kernel = "squared_exponer
mmd_sq_imq = model.compute_mmd(X_test, h = 0.5, kernel = "inverse_multi
results[f"histogram_bin_size_{bin_size}"] = (mmd_sq_se, mmd_sq_imq)
avg_bin_sizes.append(np.mean(model.bin_size))

# Plotting
ax = axs.flatten()[bin_sizes.index(bin_size)]
X_samples = model.sample(n_samples = 2000)
ax.scatter(X_train[:, 0], X_train[:, 1], s=10, label = "Training Data")
ax.scatter(X_samples[:, 0], X_samples[:, 1], s=10, alpha = 0.4, label =
ax.set_title(f"(bin size={model.bin_size})\nMMD^2 (SE) = {mmd_sq_se:.4f}
ax.legend()

fig.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()


```

```

/var/folders/dr/34chm7gn07vdb_w184vptt640000gn/T/ipykernel_70289/2696974575.
py:98: RuntimeWarning: divide by zero encountered in log
self.log_densities[valid_mask] = np.log(

```

2D Histogram Models with Different Bin Sizes




```

In [116... # Plot MMD^2 results for 2D histogram
fig, axs = plt.subplots(1, 2, figsize=(12, 6))
# SE / IMQ
results_se = [results[key][0] for key in results]
results_imq = [results[key][1] for key in results]
axs[0].scatter(avg_bin_sizes, results_se)
axs[0].set_title("MMD^2 (Squared Exponential) for 2D Histogram Models")
axs[0].set_ylabel("MMD^2")
axs[1].scatter(avg_bin_sizes, results_imq)
axs[1].set_title("MMD^2 (Inverse Multi-Quadratic) for 2D Histogram Models")
axs[1].set_ylabel("MMD^2")

```

```
fig.tight_layout()
plt.show()
```



Comment on 2D Histogram results

Both kernels (se, imq) seem to find the same optimal solution among the selected ones at a bin_size of (0.05, 0.05) - (0.1, 0.1). Every run yields slightly different results, but consistently this range is among the best performing.

In 1d, the mathematically optimal solution calculates the Interquantile Range and determines from that the bin_size. In our case, it was of around (0.22, 0.18) which we plotted for simplicity purposes at an average of bin_size = 2. This does not seem to minimize the MMD^2 in two dimensions.

At a first glance, it appears as though bin_sizes that are too small overfit the model, so that it doesn't generalize when we use our test data.

GMM

```
In [117... # Evaluate the GMM
# Bin sizes
Ms = [1, 2, 3, 5, 10, 15, 30, 50]
hs = [0.1, 0.2, 0.3, 0.5, 1.0, 1.5, 3.0, 5.0]

results = {}
# Make a grid of plots
fig, axs = plt.subplots(2, 4, figsize=(15, 10))
fig.suptitle("Gaussian Mixture Models with Different Hyperparameters", font
for M in Ms:
    for h in hs:
        model = gaussian_mixture_model(n_terms=M, steps=100)
        model.fit(X_train)
        model.sample(2000)
        model.score_samples(X_test)
```



```

mmd_sq_se = model.compute_mmd(X_test, h=h, kernel="squared_exponenti
mmd_sq_imq = model.compute_mmd(X_test, h=h, kernel="inverse_multi_qu
results[f"{M}_{h}"] = (mmd_sq_se, mmd_sq_imq)
if h == 5.0:
    print(mmd_sq_se, mmd_sq_imq)

# Plotting (This will only plot for the last bandwidth but idc)
ax = axs.flatten()[Ms.index(M)]
X_samples = model.sample(n_samples=2000)
ax.scatter(X_train[:, 0], X_train[:, 1], s=10, label="Training Data")
ax.scatter(X_samples[:, 0], X_samples[:, 1], s=10, alpha=0.4, label="Sam
ax.set_title(f"(M={M}, h={h})\nMMD^2 (SE) = {mmd_sq_se:.4f}\nMMD^2 (IMQ)
ax.legend()

fig.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()


```

```

-0.9262088913610866 -0.9641842640662349
-0.926222711862457 -0.9641956759677022
-0.9262333258308659 -0.9642005576205309
-0.9261748661558955 -0.9641732226228716
-0.9262184295202183 -0.96419353689215
-0.9261861820689726 -0.9641779604148133
-0.9262276826082626 -0.9641980520250333
-0.9260091749126373 -0.9640921344016296

```

Gaussian Mixture Models with Different Hyperparameters



```

In [118... # Plot MMD^2 results for GMM
fig, axs = plt.subplots(1, 2, figsize=(12, 12))
# SE / IMQ
results_se = [results[key][0] for key in results]
results_imq = [results[key][1] for key in results]
Ms_list = [int(key.split("_")[1]) for key in results]
hs_list = [float(key.split("_")[2]) for key in results]


```



```

scatter1 = axs[0].scatter(Ms_list, results_se, c=hs_list, cmap='viridis')
axs[0].set_title("MMD^2 (Squared Exponential) for GMMs")
axs[0].set_xlabel("Number of Mixture Components (M)")
axs[0].set_ylabel("MMD^2")
cbar1 = fig.colorbar(scatter1, ax=axs[0])
cbar1.set_label('Bandwidth (h)')
scatter2 = axs[1].scatter(Ms_list, results_imq, c=hs_list, cmap='viridis')
axs[1].set_title("MMD^2 (Inverse Multi-Quadratic) for GMMs")
axs[1].set_xlabel("Number of Mixture Components (M)")
axs[1].set_ylabel("MMD^2")
cbar2 = fig.colorbar(scatter2, ax=axs[1])
cbar2.set_label('Bandwidth (h)')
fig.tight_layout()
plt.show()

```



Comment on GMM results

There are now multiple dimensions to this, because we test the model for different numbers of components and different bandwidths. We first comment on the total

number of components, where at some point all models seem to converge. However, the models with larger bandwidth converge much faster which is totally weird because obviously for a small number of Gaussian terms, the model does not generalize well. Maybe a mistake in our calculation.

2. Higher dimensional data

```
In [119... # Import dataset
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split

digits = load_digits()
X, y = digits.data, digits.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```

```
In [120... # Import Models
from sklearn.ensemble import RandomForestClassifier # In addition
from sklearn.neighbors import KDTree
from sklearn.mixture import GaussianMixture
from sklearn.neighbors import KernelDensity
```

```
In [121... # Fit the models
# KDTree (instead of 2d Histogram)
kdtree = KDTree(X_train)
# GMM
gmm_10 = GaussianMixture(n_components=10, random_state=42)
gmm_10.fit(X_train)
gmm_100 = GaussianMixture(n_components=100, random_state=42)
gmm_100.fit(X_train)
# KDE
kde = KernelDensity(kernel='gaussian', bandwidth=1.0)
kde.fit(X_train)
# Random Forest
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
```

```
Out[121... ▼ RandomForestClassifier ⓘ ⓘ
  ► Parameters
```

```
In [122... # Sample
n_samples = 1000
# KDTree Sampling
dist, indices = kdtree.query(X_train[np.random.choice(X_train.shape[0], n_sa
kdtree_samples = X_train[indices.flatten()]
# GMM Sampling
gmm_10_samples, _ = gmm_10.sample(n_samples)
gmm_100_samples, _ = gmm_100.sample(n_samples)
# KDE Sampling
```

```
kde_samples = kde.sample(n_samples)
# Random Forest Sampling
rf_samples = rf.apply(X_train[np.random.choice(X_train.shape[0], n_samples)])
rf_samples = X_train[rf_samples[:, 0]] # Using the first tree for sampling
```

```
In [123... # Compute MMD^2
def compute_mmd(X, h = 1, kernel = "squared_exponential"):
    # Implement MMD^2 metric with se, imq kernels
    # Need to vectorize MMD calculation because for loops (sums) are very in
    # MMD^2 = Term1 + Term2 - Term3, can neglect Term1 because it is model i
    """
    MMD^2 = 0
    for i in range(1,N):
        for i_prime != i:
            MMD^2 += k(x_i, x_i_prime)

    for j in range(1, N_prime):
        for j_prime != j:
            MMD^2 += k(x_prime_j, x_prime_j_prime)

    for i in range(1, N):
        for j in range(1, N_prime):
            MMD^2 += -2 k(x_i, x_prime_j)
    """
    # Instead, we want to calculate it using something like this:
    """
    MMD^2 = k(distance_matrix_x).sum()
            + k(distance_matrix_x_prime).sum()
            - 2 k(distance_matrix_xx_prime).sum()
    """
    # Calculate Distance matrices D as:
    # D[i, j] = || x_i - x_i_prime ||_2^2 = ||x||_2^2 + ||x_prime||_2^2 - 2
    # D_XX = self.calculate_distance_matrix(X, X) Can be neglected
    D_XprimeXprime = calculate_distance_matrix(samples, samples)
    D_XXprime = calculate_distance_matrix(X, samples)

    # Depending on selected kernel, calculate MMD^2
    N = X.shape[0]
    Nprime = samples.shape[0]
    if kernel == "squared_exponential":
        MMD_sq = (squared_exponential(D_XprimeXprime, h).sum() / (Nprime **
            - 2 * squared_exponential(D_XXprime, h).sum() / (N * Npr
    elif kernel == "inverse_multi_quadratic":
        MMD_sq = (inverse_multi_quadratic(D_XprimeXprime, h).sum() / (Nprime
            - 2 * inverse_multi_quadratic(D_XXprime, h).sum() / (N *
    else:
        raise ValueError("Please select either 'squared_exponential' or 'inv
    """
    In the lecture we derived the formula where we sum over i, then sum over
    k(i, i) are not 0, but 1. We still sum over them in the vectorization. F
    we adjust the denominator to N^2, N'^2 instead of N(N-1), N'(N'-1)
    """

    return MMD_sq

def calculate_distance_matrix(A, B):
```

```

#  $D_{ij} = ||A_i - B_j||^2 = ||A_i||^2 + ||B_j||^2 - 2 A_i @ B_j.T$ 
#  $A_{sq}$  in  $(N_A, 1)$ 
A_sq = np.sum(A ** 2, axis = 1, keepdims = True)
#  $B_{sq}$  in  $(1, N_B)$ 
B_sq = np.sum(B ** 2, axis = 1, keepdims = True).T # .T otherwise won't
AB_T = A @ B.T
D_AB = A_sq + B_sq - 2 * AB_T
return D_AB

def squared_exponential(D_AB, h):
    return exp(- D_AB / (h ** 2))

def inverse_multi_quadratic(D_AB, h):
    return 1 / sqrt(1 + D_AB / (h ** 2))

```

```

In [131... # Compute MMD^2 for each model
models_samples = {
    "KDTree": kdtree_samples,
    "GMM_10": gmm_10_samples,
    "GMM_100": gmm_100_samples,
    "KDE": kde_samples
}

mmd_results = {}
for model_name, samples in models_samples.items():
    mmd_sq_se = compute_mmd(X_test, h = 1.0, kernel = "squared_exponential")
    mmd_sq_imq = compute_mmd(X_test, h = 1.0, kernel = "inverse_multi_quadra")
    mmd_results[model_name] = (mmd_sq_se, mmd_sq_imq)

print("MMD^2 Results for Different Models:")
for model_name, (mmd_se, mmd_imq) in mmd_results.items():
    print(f"{model_name}: MMD^2 (SE) = {mmd_se:.4f}, MMD^2 (IMQ) = {mmd_imq:.4f}")

```

```

MMD^2 Results for Different Models:
KDTree: MMD^2 (SE) = 0.0017, MMD^2 (IMQ) = -0.0198
GMM_10: MMD^2 (SE) = 0.0010, MMD^2 (IMQ) = -0.0204
GMM_100: MMD^2 (SE) = 0.0010, MMD^2 (IMQ) = -0.0205
KDE: MMD^2 (SE) = 0.0010, MMD^2 (IMQ) = -0.0205


```

```

In [132... # Plot examples
fig, axs = plt.subplots(5, 4, figsize=(12, 6))
for i in range(4):
    model_name = list(models_samples.keys())[i]
    samples = models_samples[model_name]
    for j in range(5):
        axs[j, i].imshow(samples[j].reshape(8, 8), cmap='gray')

        if j == 0:
            axs[0, i].set_title(f"{model_name}")

```




```
In [133... # Check model accuracy with RF
# Print score, accuracy
print("Random Forest Classifier Accuracy on Test Set:", rf.score(X_test, y_t
```

Random Forest Classifier Accuracy on Test Set: 0.9722222222222222

```
In [134... # Predict labels from data generated by models with RF
y_hats_models = {
    model_name: rf.predict(samples) for model_name, samples in models_sample
}

# Show if uniformly distributed
y_hats_counts = {
    model_name: np.bincount(y_hats, minlength = 10) for model_name, y_hats i
}

# Make some bar charts
fig, axs = plt.subplots(1, 4, figsize=(20, 5))
for i in range(4):
    model_name = list(y_hats_counts.keys())[i]
    counts = y_hats_counts[model_name]
    axs[i].bar(range(10), counts)
    axs[i].set_title(f"Predicted Distribution for {model_name}")
    axs[i].set_xlabel("Digit Label")
    axs[i].set_ylabel("Count")
plt.show()
```



```
In [136... # Perform hypothesis test on each distribution
from scipy.stats import chisquare
alpha = 0.05
for model_name, counts in y_hats_counts.items():
    chi2_stat, p_value = chisquare(counts)
    print(f"{model_name}: Chi2 Stat = {chi2_stat:.4f}, p-value = {p_value:.4f}")
    if p_value < alpha:
        print(f" Reject null hypothesis for {model_name}")
    else:
        print(f" Fail to reject null hypothesis for {model_name}.")
```

KDTree: Chi2 Stat = 16.1200, p-value = 0.0644

Fail to reject null hypothesis for KDTree.

GMM_10: Chi2 Stat = 28.5600, p-value = 0.0008

Reject null hypothesis for GMM_10

GMM_100: Chi2 Stat = 7.3200, p-value = 0.6038

Fail to reject null hypothesis for GMM_100.

KDE: Chi2 Stat = 8.9600, p-value = 0.4410

Fail to reject null hypothesis for KDE.

Comment on results

Upon visual inspection, it is clear that the GMM model with only 10 Gaussian terms does not generalize well to the data and it returns very uneven results, having nearly twice as many "0" labels as "1" labels. The GMM model with 100 Gaussian terms fairs much better in comparison and the KDE distribution is also fairly equal. Visually, it is not clear whether the KDTree distribution passes the vibe check, but a simple χ^2 test yields that we can't reject the null hypothesis, although based on the p-value it is just slightly larger than our alpha value. Likely in a different run, the result could fall on either side.

In []:

In []: